

## Project 3

Robin Lin

2022-11-04

```
## The author of this project is Sihong Lin, which can also be referred to as  
## Robin Lin.
```

```
## The student number of the author is s2435943.
```

```
## The author declares that everything in the project is his own work, and he  
## promises that he did not plagiarise from others.
```

```
## The url of the GitHub repository is listed as follows.  
## https://github.com/RobinLam2435943/Statistical-Programming-Individual-  
Projects.git
```

```
## -----  
---
```

```
## Code to smooth with basis expansions and penalties.
```

```
## The y values consist of a smoothing function of x, as well as the error  
terms.
```

```
## The smoothing function is a linear combination of k evenly-spaced B-  
splined  
## basis functions, and the coefficients are to be estimated.
```

```
## In order to avoid over-fitting, some smoothing penalty is imposed, making  
each  
## coefficient vary smoothly from its neighbouring ones.
```

```
## The model is hence estimated by the penalised least squares, and the  
smoothing  
## parameter is chosen in order to minimise the generalised cross validation  
## (GCV) criterion.
```

```
## This project aims at fitting P-splines to the x, y data, and selecting  
## the optimal smoothing parameter based on GCV criterion. Also, some details  
of  
## the output of the method function are needed. Predictions are to be made,  
## and some diagnostic plots are to be sketched.
```

```
pspline <- function(x, y, k = 20, logsp = c(-5, 5), bord = 3, pord = 2, ngrid  
= 100){
```

*## This function aims at performing the best fit to the smoothing function.  
The*

*## output is a bunch of lists containing detailed information of the model.*

*## x, y are data points, and k is the number of evenly-spaced B-splined  
basis*

*## functions. 'logsp' represents the boundaries of smoothing functions in  
log*

*## scale. 'bord' is the B-spline order, and 'pord' is the penalty order of  
## difference. 'ngrid' is the number of smoothing parameters to try.*

*dk <- diff(range(x)) / (k - bord) # Knot spacing.*

*knots <- seq(min(x) - dk \* bord, by = dk, length = k + bord + 1) # Knots  
# generating.*

*X <- splines::splineDesign(knots, x, ord = bord + 1, outer.ok = TRUE) #  
Data*

*# matrix.*

*D <- diff(diag(k), differences = pord) # Penalisation matrix.*

*qrX <- qr(X) # Obtains QR-factorisation of X, i.e.,  $X = QR$ .*

*Q <- qr.Q(qrX) # Matrix Q.*

*R <- qr.R(qrX) # Matrix R.*

*eig <- eigen(t(solve(R)) %\*% t(D) %\*% D %\*% solve(R)) # Obtains eigen-  
# decomposition of  $(D * R ^ (-1)) ^ T * (D * R ^ (-1))$ , i.e.,  
#  $(D * R ^ (-1)) ^ T * (D * R ^ (-1)) = U * \Lambda * (U ^ T)$ .*

*U <- eig\$vectors # Matrix U.*

*Lam <- diag(eig\$values) # Matrix  $\Lambda$ .*

*logSP <- seq(from = logsp[1], to = logsp[2], length.out = ngrid) # Equally  
# separated smoothing parameters in log scale.*

*SP <- exp(logSP) # Smoothing parameters.*

*gcv <- function(sm\_par){*

*## This function aims at reckoning GCV criterion given a smoothing  
parameter.*

*revised\_Lam <- diag(diag(1 + sm\_par \* Lam)) # Matrix  $(I + \text{smoothing\_par} *  
# \Lambda)$ .*

*A <- solve(R) %\*% U %\*% solve(revised\_Lam) %\*% t(U) # Matrix  $(R ^ (-1) *  
U *$*

*# revised\_Lam ^ (-1) \* U ^ T).*

*coef <- A %\*% (t(Q) %\*% y) # Estimates of coefficients.*

*fitted <- X %\*% coef # Fitted Values.*

*edf <- sum(diag(solve(revised\_Lam))) # Effective degrees of freedom,  
reckoned*

```

    # by taking the trace on the inverse of 'revised_Lam'.
    sig2 <- sum((y - fitted) ^ 2) / (nrow(X) - edf) # Residual Variance.
    gcv <- sig2 / (nrow(X) - edf) # GCV criterion.
    return(gcv)
}

GCV <- as.numeric(lapply(X = SP, FUN = gcv)) # Stores all the GCV values.
sp <- SP[which.min(GCV)] # Selects the smoothing parameter with the
# smallest GCV criterion.

revised_Lam <- diag(diag(1 + sp * Lam)) # Matrix (I + smoothing_par *
# \Lambda).
A <- solve(R) %*% U %*% solve(revised_Lam) %*% t(U) # Matrix (R ^ (-1) * U
*
# revised_Lam ^ (-1) * U ^ T).
coef <- A %*% (t(Q) %*% y) # Estimates of coefficients.

fitted <- X %*% coef # Fitted Values.
edf <- sum(diag(solve(revised_Lam))) # Effective degrees of freedom,
reckoned
# by taking the trace on the inverse of 'revised_Lam'.
sig2 <- sum((y - fitted) ^ 2) / (nrow(X) - edf) # Residual variance.
V <- sig2 * A %*% t(solve(R)) # Covariance matrix for the coefficients.
r2 <- 1 - (nrow(X) - 1) * sig2 / sum((y - mean(y)) ^ 2) # Model R-squared.
gcv <- sig2 / (nrow(X) - edf) # GCV criterion.

newlist <- list('x' = x, 'y' = y, 'B_spline_Order' = bord,
'Penalty_Order_of_Difference' = pord, 'Number_of_Basis_Functions' = k,
'Knots' = knots, 'Smoothing_Parameter' = sp, 'Coefficients' = coef,
'Fitted_Values' = fitted, 'Effective_Degrees_of_Freedom' = edf,
'Residual_Variance' = sig2, 'Residual_Std' = sqrt(sig2),
'Cov_for_Coefficients' = V, 'R_Squared' = r2, 'Generalised_Cross_Validation'
= gcv)
# Creates a list containing a bunch of details of the model.

return(newlist)
}

print.pspline <- function(m){

  ## This function aims at showing the EDF, k, residual std, r ^ 2, and GCV
of
  ## the model, given the method function 'pspline'.

  cat('Order', m$B_spline_Order, 'p-spline with order',
m$Penalty_Order_of_Difference, 'penalty', '\n')

  cat('Effective degrees of freedom:', m$Effective_Degrees_of_Freedom,
'Coefficients:', m$Number_of_Basis_Functions, '\n')

```

```

cat('residual std dev:', m$Residual_Std, 'r-squared:', m$R_Squared, 'GCV:',
m$Generalised_Cross_Validation, '\n')

newlist <- list('gcv' = m$Generalised_Cross_Validation, 'edf' =
m$Effective_Degrees_of_Freedom, 'r2' = m$R_Squared)
# Stores the values of gcv, edf, and r2.

invisible(newlist) # Silently returns the new list.
}

predict.pspline <- function(m, x, se = TRUE){

  ## This function aims at making predictions from the model given new x
  values
  ## within the range of the original data, as well as providing the standard
  ## errors of the model.

  Xp <- splines::splineDesign(m$Knots, x, ord = m$B_spline_Order + 1,
outer.ok = TRUE)
  # Data matrix.

  coef <- m$Coefficients # Estimates of coefficients.
  fitted <- Xp %*% coef # Fitted Values.

  if(se){ # Standard errors and fitted values are to be returned.
    V <- m$Cov_for_Coefficients # Covariance matrix for the coefficients.
    std_error <- rowSums(Xp * (Xp %*% V)) ^ .5 # Standard Errors.
    newlist <- list('fit' = fitted, 'se' = std_error) # Stores the fitted
  values
    # and the standard errors.
  }else{ # Only fitted values are to be returned.
    newlist <- list('fit' = fitted) # Stores only the values of fitted
  values.
  }

  return(newlist)
}

plot.pspline <- function(m){

  ## This function aims at sketching 3 diagnostic plots of the model.

  ## The first one is the plot of the original data set, with the smoothing
  ## line being overlaid, along with the 95% credible intervals for the
  smooth.
  ## In order to obtain the 95% credible intervals, both the lower and the
  ## upper limits should be reckoned. These two limits, along with their
  ## corresponding x values, should be returned silently.

```

*## The second one is the plot of the residuals against fitted values.*

*## The third one is the Normal Q-Q plot.*

```
CL <- .95 # Credible Level.
```

```
CV <- qnorm((1 - CL) / 2, lower.tail = FALSE) # Critical Value.
```

```
ll <- m$Fitted - CV * predict.pspline(m, m$x)$se # Lower Confidence Limit.
```

```
ul <- m$Fitted + CV * predict.pspline(m, m$x)$se # Upper Confidence Limit.
```

*## The First Plot*

```
plot(m$y ~ m$x, main = 'Smoothing Plot with 95% Credible Intervals', xlab =  
'x', ylab = 'y', type = 'p', col = 'white')  
# Sketches an empty graph.
```

```
lines(ll ~ m$x, col = 'blue', lwd = 3) # Sketches a line of lower  
confidence  
# limit.
```

```
lines(ul ~ m$x, col = 'blue', lwd = 3) # Sketches a line of upper  
confidence  
# limit.
```

```
polygon(c(m$x, rev(m$x)), c(ll, rev(ul)), col = 'grey', border = NA)
```

```
# Sketches a grey-shaded area representing the 95% credible intervals.
```

```
lines(m$Fitted ~ m$x, col = 'red', lwd = 3) # Sketches the smoothing line.
```

```
points(m$y ~ m$x) # Sketches the data points.
```

```
legend('bottomright', legend = c('Smoothing Curve', '95% Credible Interval  
Bounds'), lwd = 3, col = c('red', 'blue'), cex = .7) # Introduces the  
legends.
```

*## The Second Plot*

```
plot((m$y - m$Fitted_Values) ~ m$Fitted_Values, xlab = 'Fitted Values',  
ylab = 'Residuals', main = 'Residual vs. Fitted Values')
```

```
# sketches the 'Residuals vs. Fitted' plot.
```

```
abline(lm((m$y - m$Fitted_Values) ~ m$Fitted_Values), col = 'red', lwd = 3)  
# Adds a line representing the trend.
```

*## The Third Plot*

```
qqnorm(m$y - m$Fitted_Values) # Normal Q-Q plot.
```

```
qqline(m$y - m$Fitted_Values, col = 'red', lwd = 3) # Normal Q-Q line.
```

```
newlist <- list('ll' = ll, 'ul' = ul, 'x' = m$x) # Stores the lower limits,  
# the upper limits, as well as the corresponding x values to the new list.  
invisible(newlist) # Silently returns the new list.
```

```
}
```

```

## Load the data set.
library(MASS)
x <- mcycle$times
y <- mcycle$accel

## Apply the 'pspline' function.
m <- pspline(x, y)

## Print EDF, k, residual std, r ^ 2, and GCV of the model.
print.pspline(m)

## Order 3 p-spline with order 2 penalty
## Effective degrees of freedom: 11.22137 Coefficients: 20
## residual std dev: 22.67567 r-squared: 0.7797937 GCV: 4.222302

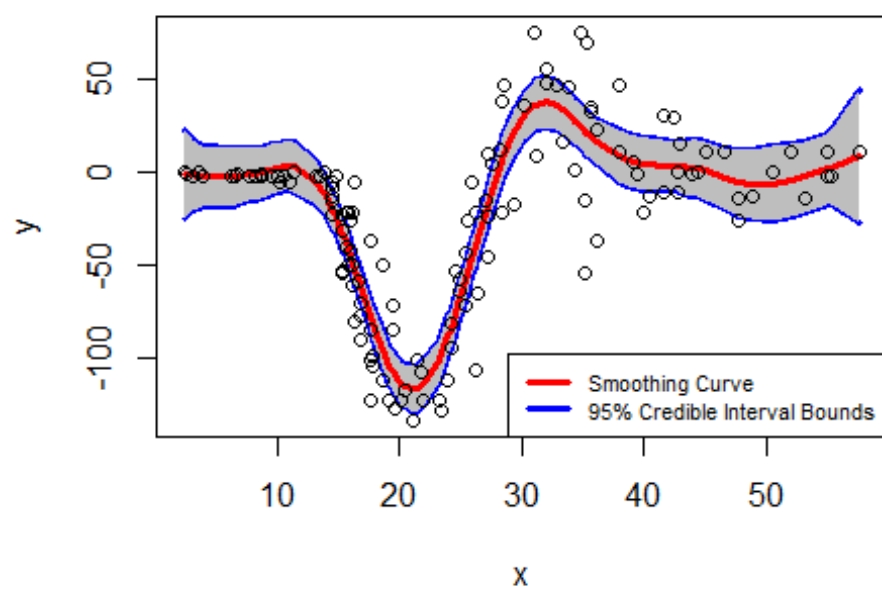
## Make predictions.
new_x<- x[1 : 10]
predict.pspline(m, new_x)

## $fit
##           [,1]
## [1,] -1.143021
## [2,] -1.318180
## [3,] -1.819277
## [4,] -2.122588
## [5,] -2.391554
## [6,] -2.885941
## [7,] -2.736742
## [8,] -2.630237
## [9,] -1.760593
## [10,] -1.246979
##
## $se
## [1] 12.369838 11.501348  9.580341  8.831780  8.449088  8.298085  8.019304
## [8]  7.864015  7.222621  7.092230


## Sketch the diagnostic plots.
plot.pspline(m)

```

### Smoothing Plot with 95% Credible Intervals



### Residual vs. Fitted Values



**Normal Q-Q Plot**

